

# Hive – A Petabyte Scale Data Warehouse Using Hadoop

Ashish Thusoo, Joydeep Sen Sarma, Namit Jain, Zheng Shao, Prasad Chakka, Ning Zhang, Suresh Antony, Hao Liu and Raghatham Murthy

*Facebook Data Infrastructure Team*

**Jose Abreu**  
May 9<sup>th</sup>

## **A Comparison of Approaches to Large-Scale Data Analysis**

Andrew Pavlo  
Brown University  
[pavlo@cs.brown.edu](mailto:pavlo@cs.brown.edu)

Erik Paulson  
University of Wisconsin  
[epaulson@cs.wisc.edu](mailto:epaulson@cs.wisc.edu)

Alexander Rasin  
Brown University  
[alexr@cs.brown.edu](mailto:alexr@cs.brown.edu)

Daniel J. Abadi Yale University <a href="mailto:dna@cs.yale.edu">dna@cs.yale.edu</a>	David J. DeWitt Microsoft Inc. <a href="mailto:dewitt@microsoft.com">dewitt@microsoft.com</a>	Samuel Madden M.I.T. CSAIL <a href="mailto:madden@csail.mit.edu">madden@csail.mit.edu</a>	Michael Stonebraker M.I.T. CSAIL <a href="mailto:stonebraker@csail.mit.edu">stonebraker@csail.mit.edu</a>
--	---	---	---

# Hive

- Open-source data warehousing solution built on top of Hadoop.
- Supports queries written in HiveQL. (Derived from SQL.)
- Used to reduce the complex nature of writing Map-Reduce programs on Hadoop.
- Used to address requirements at Facebook. Ranging from the Ad Network to Facebook's Lexicon product.
- Supports hundreds of terabytes of data.
- Allows for legacy data to be incorporated into a table without transforming data.

# Implementation

- Data is stored in tables of rows and columns. Each column has a associated type (Primitive Type or Complex Type)
- “Complex types are templated and can be composed to generate types of arbitrary complexity.”
- LazySerDe – Deserializes rows into internal objects lazily so cost of deserialization of a column is incurred only if the column of the row is needed in query expression.
- Execution Engine - Each dependent task is only executed if all of its prerequisites have been executed.

# Analysis

- Being Open-Source is always a plus.
- Being in a familiar language helps those that are coming into it.
- All or nothing task scheduling.
- Being able to transfer legacy tables means that you don't need to rip out everything if you decide to move over to Hive.
- Do not like the fact that you are not able to insert or update a table without overwriting the entire table.

# Comparisons to “A Comparison of Approaches to Large-Scale Data Analysis”

- Parallel systems seem to perform better than MapReduce systems.
- Serialized data formats and compression features of Hadoop systems are much slower at data loading.
- MapReduce has more overhead than Parallel systems.
- Database systems were lacking during tests.
- MapReduce needs 1000 nodes to make 100 nodes on a Parallel System.

# Advantages and Disadvantages

- Hadoop systems are easier to implement.
- No need for custom input handlers.
- Hadoop is the most popular open source implementation.
- Much slower than parallel systems.
- “Hive currently does not support inserting into an existing table or data partition and all inserts overwrite the existing data.”
- Speed is limited with how fast each table can be read off a disk. (Parallel can use clustered indexes to do this faster.)